CS378: Final Project - Data Artifacts
https://github.com/sarimaleem/cs378_fp

Sarim Aleem
ska2222

Abstract

TODO: write an abstract describing the core motivation and findings from your project

1 Introduction (5pt)

TODO: This section should include three paragraphs. The first paragraph is to briefly describe task and data. The second paragraph should describe the results of your analysis, and the third paragraph describing your fix and main experimental take aways.

The model that was used was [1]

2 Task/Dataset/Model Description (15pt)

TODO: Describe the task/dataset/model you are working on. Clearly define your task with mathematical notations. Describe your learning algorithm. You must formally specify the loss function, stopping criteria, training data, etc used for the model you are analyzing in the next section. Remember, every notation you use must be defined.

3 Performance Analysis (25pt)

The model was trained for 3 iterations on the SNLI training dataset. Each iteration has 550152 training examples, the batch size for the training examples was 32. All training was done on google colab in the cloud, and training took about 25 minutes.

Overall Accuracy Statistics

An initial evaluation of the model shows that it has an accuracy of 88%. We also did further evaluation to show statistics about which labels it misclassified most using a confusion matrix.

Table 1: Development Set Evaluation Metrics

loss	0.315
accuracy	0.8865
runtime	20.2629
examples/s	485.715
steps/s	60.751

Model Metrics for 9842 examples

Table 2: Confusion Matrix

	Predicted				
True		*0*	*1*	*2*	**Total**
	0	2996	252	81	3329
	1	213	2787	235	3235
	2	79	257	2942	3278
	Total	3288	3296	3258	**9842**

0=Entailment, 1=Neutral, 2=Contradiction

Table 3: Percent mispredictions

	Predicted			
True		*0*	*1*	*2*
	0	89.99	7.78	2.47
	1	6.39	86.15	7.16
	2	2.37	7.94	89.74

0=Entailment, 1=Neutral, 2=Contradiction

3.1 Confusion Matrix Analysis

In general, the model seems to do well with Separating Entailment and Contradiction. For example, it only incorrectly identifies 2.47% of entailment examples as contradiction and 2.37% of contradiction examples as entailement.

However, the model struggles more to understand the relationship between, entailment and neutral, as well as contradiction and neutral. For example, it mispredicted over 7% of entailments and contradictions as neutrals. As long as mispredicting over 6% of neutrals as entailments and over 7% of of neutrals as contradictions.

3.2 Manual Pattern Analysis

It's difficult to say how the model is evaluating sentences, since it's impossible to fully understand its weights. Nevertheless, there do seem to be some patterns one evaluating mistakes that the model makes. One pattern that seems to happen is that the model seems to be unable to understand different words that mean the same thing in a context. For example, the following is an error that the model made.

In the example with hypothesis *A man and a woman are looking at produce on display* and premise *A man and woman are staring at heads of lettuce.* The model is unable to distinguish that lettuce is a type of produce and instead makes assumption that they are different, thus leading it to think they are contradictions.

3.3 Using Contast sets to evaluate the Model

subsectionContrast Sets

In their article *Evaluating models' local decision boundaries via contrast sets* [2], Gardner et al. explain the concept of a *contrast set.* A contrast set is a small but meaninful alteration in the input data that typically leads to a different gold label.

In order to test if the model is finding artifacts in the data or genuinely evaluating it, we annotated 35 different examples in the validation data that the model predicted correctly, and then tested to see if the model was able to predict them correctly again.

The results of the model showed a significant dip in accuracy, with accuracy going down to 77% from 89%. The model mispredicted 8/35 examples. One pattern that emerged from the mispredictions is that the model fails to see correlations between various words and also fails to understand the relationships between various words. The following is an example of this phenomenon:

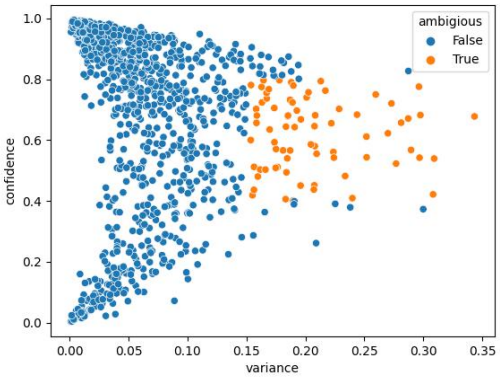

Figure 1: Data Cartogram of 2000 training examples

> premise: Two men on vehicles competing in a race.
> hypothesis: Men are riding cars on the street.
> label: neutral
> predicted: contradiction

The neural network is essentially unable to distinguish that different words can have the same meaning (in this case vehicles and cars).

4 Describing Your Fix (20pt)

It is clear that the issue with the model is that it fails to understand ambiguous training examples. This is especially seen given that the model largely confuses, as seen earlier, neutral and entailment and instances of contradiction.

4.1 Dataset Cartography

Training on labels that are ambiguous, that have high variablity, should in theory help with this issue, even more so than hard to learn labels. An ambiguous label is defined as a label that has high variance when predicted by different iterations of the model. For example, a model that is trained on 500 iterations would predict the input differently than a model trained with 1000 iterations or a model trained with 1500 iterations. Furthermore, the average confidence of that training example is not high or low. [3] We reproduced the

Table 4: Examples of model misunderstanding words, or not considering their importance

premise	Hypothesis	Gold	Predicted
People are throwing tomatoes at each other	The people are having a food fight	entailment	contradiction
A man and a woman are looking at produce on display.	A man and women are staring at heads of lettuce.	neutral	contradiction
Two men sitting on a subway are reading, with coats and scarves on, but have seemed to have lost their pants.	The men are wearing pants.	contradiction	entailment
Two men are in an electronics workshop, working on computers or equipment	The men are unaware of what computers are.	contradiction	neutral

code to analyze data, and have plotted 500 training examples in figure 1.

In order to deal with this, we first created a program that finds ambiguous examples. We found the variance of the training example based on its confidence of the gold label from a modeled trained on 500 iterations, 1000 iterations, 1500 iterations, 2000 iterations, 2500 iterations, and 3000 iterations. If th variance was greater than 0.2, and the confidence of the model was in between 0.6 and 0.7, the example was deemed ambiguous. According to *Swayamdipta et al.*, training solely ambiguous examples leads to to greater performance than training on the whole dataset, only easy to learn examples, or only hard to learn examples. Thus it is logical to train on these examples.

We pulled 1655 ambiguous samples from the training data and trained on those examples.

5 Evaluating Your Fix (25pt)

TODO: Your writeup should address how effective is your fix, how broadly applicable is your fix, etc. Providing a single number (overall accuracy) is necessary but not sufficient here. For example, if your change made the model better on challenging NLI examples, you could try to quantify that on one or more slices of the data, give examples of predictions that are fixed, or even use model interpretation techniques to try to support claims about how your improved model is doing its "reasoning." (You can look at the papers listed above to get a sense of how to do such fine-grained evaluation). You should report results from a baseline approach (your initial trained model) as well as your "best" method. If doing your own project, baselines such as majority class, random, or a linear classifier are important to see. **Ablations**: If you tried several things, analyze the contribution from each one. These should be *minimal* changes to the same system; try running things with just one aspect different in order to assess how important that aspect is. This part of the report should be at least one page.

6 Related Work (5pt)

TODO: Briefly discuss prior research papers related to your approach. This will likely some papers in the project description document. How is your approach different from existing studies?

7 Conclusion (5pt)

TODO: Brief conclusion summarizing findings from both numerical results and qualitative analysis.

(Optional) AI Assistance

TODO: If you have used any AI toolkit (either for writing assistance or code assistance) for your final project, please describe it here.

References

[1] Kevin Clark, Minh-Thang Luong, Quoc V. Le, and Christopher D. Manning. Electra: Pre-

training text encoders as discriminators rather than generators, 2020.

[2] Matt Gardner, Yoav Artzi, Victoria Basmova, Jonathan Berant, Ben Bogin, Sihao Chen, Pradeep Dasigi, Dheeru Dua, Yanai Elazar, Ananth Gottumukkala, et al. Evaluating models' local decision boundaries via contrast sets. *arXiv preprint arXiv:2004.02709*, 2020.

[3] Swabha Swayamdipta, Roy Schwartz, Nicholas Lourie, Yizhong Wang, Hannaneh Hajishirzi, Noah A. Smith, and Yejin Choi. Dataset cartography: Mapping and diagnosing datasets with training dynamics, 2020.